TransNets: Learning to Transform for Recommendation

Rose Catherine   **William Cohen**

School of Computer Science
Carnegie Mellon University
`{rosecatherinek,wcohen}@cs.cmu.edu`

ABSTRACT

Recently, deep learning methods have been shown to improve the performance of recommender systems over traditional methods, especially when review text is available. For example, a recent model, *DeepCoNN*, uses neural nets to learn one latent representation for the text of all reviews written by a target user, and a second latent representation for the text of all reviews for a target item, and then combines these latent representations to obtain state-of-the-art performance on recommendation tasks. We show that (unsurprisingly) much of the predictive value of review text comes from reviews of the target user for the target item. We then introduce a way in which this information can be used in recommendation, even when the target user's review for the target item is not available. Our model, called *TransNets*, extends the *DeepCoNN* model by introducing an additional latent layer representing the target user-target item pair. We then regularize this layer, at training time, to be similar to another latent representation of the target user's review of the target item. We show that TransNets and extensions of it improve substantially over the previous state-of-the-art.

1 INTRODUCTION

Using review text for predicting ratings has been shown to greatly improve the performance of recommender systems [4, 22, 24], compared to Collaborative Filtering (CF) techniques that use only past ratings [18, 33]. Recent advances in Deep Learning research have made it possible to use Neural Networks in a multitude of domains including recommender systems, with impressive results. Most neural recommender models [3, 10, 16, 21, 40] have focussed on the *content* associated with the user and the item, which are used to construct their latent representations. Content associated with a user could include their demographic information, socioeconomic characteristics, their product preferences and the like. Content linked to an item could include their price, appearance, usability and similar attributes in the case of products, food quality, ambience, service and wait times in the case of restaurants, or actors, director, genre, and similar metadata in the case of movies. These representations are then fed into a CF-style architecture or a regression model to make the rating prediction.

Review text, unlike content, is not a property of only the user or only the item; it is a property associated with their joint interaction. In that sense, it is a *context* [1] feature. Only a few neural net models [2, 34, 44] have been proposed to date that use review text for predicting the rating. Of these, the most recent model, *Deep Cooperative Neural Networks* (DeepCoNN) [44] uses neural nets to learn a latent representation for the user from the text of all reviews written by her and a second latent representation for the item from the text of all reviews that were written for it, and then combines these two representations in a regression layer to obtain state-of-the-art performance on rating prediction. However, as we will show, much of the predictive value of review text comes from
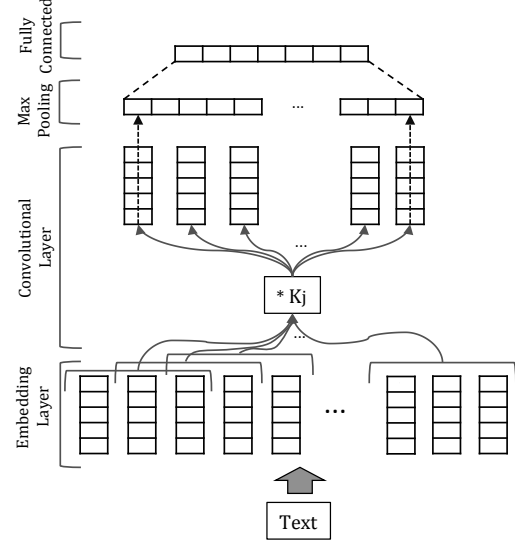

Figure 1: The *CNN Text Processor* architecture

reviews of the target user for the target item, which can be assumed to be available only at training time, and is not available at test time. In this paper, we introduce a way in which this information can be used in training the recommender system, such that when the target user's review for the target item is not available at the time of prediction, an approximation for it is generated, which is then used for predicting the rating. Our model, called **Trans**formational **Net**works (TransNets), extends the DeepCoNN model by introducing an additional latent layer representing an approximation of the review corresponding to the target user-target item pair. We then regularize this layer, at training time, to be similar to the latent representation of the actual review written by the target user for the target item. Our experiments illustrate that TransNets and its extensions give substantial improvements in rating prediction.

The rest of this paper is organized as follows. The proposed model and architecture are discussed in detail in Section 2. The experiments and results are discussed in Section 3. Section 4 summarizes related work, and we conclude in Section 5.

2 PROPOSED METHOD

2.1 CNNs to process text

We process text using the same approach as the current state-of-the-art method for rating prediction, *DeepCoNN* [44]. The basic building block, referred to as a *CNN Text Processor* in the rest of this paper, is a Convolutional Neural Network (CNN) [20] that inputs a sequence of words and outputs a n-dimensional vector representation for the input, i.e., the *CNN Text Processor* is a function $\Gamma : [w_1, w_2, ..., w_T] \rightarrow \mathbb{R}^n$. Figure 1 gives the architecture of the *CNN Text Processor*. In the first layer, a word embedding function
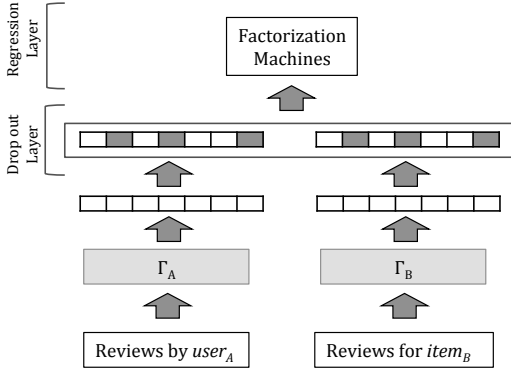
Figure 2: DeepCoNN model for predicting rating

$f : M \to \mathbb{R}^d$ maps each word in the review that are also in its M-sized vocabulary into a d dimensional vector. The embedding can be any pre-trained embedding like those trained on the GoogleNews corpus using *word2vec*[1][27], or on Wikipedia using GloVe[2] [30]. These word vectors are held fixed throughout the training process.

Following the embedding layer is the Convolutional Layer, adapted to text processing [8]. It consists of m neurons each associated with a filter $K \in \mathbb{R}^{t \times d}$, where t is a window size, typically $2 - 5$. The filter processes t-length windows of d-dimensional vectors to produce features. Let $V_{1:T}$ be the embedded matrix corresponding to the T-length input text. Then, j^{th} neuron produces its features as:

$$z_j = \alpha(V_{1:T} * K_j + b_j)$$

where, b_j is its bias, $*$ is the convolution operation and α is a non-linearity like Rectified Linear Unit (ReLU) [28] or $tanh$.

Let $z_j^1, z_j^2, ... z_j^{(T-t+1)}$ be the features produced by the j^{th} neuron on the sliding windows over the embedded text. Then, the final feature corresponding to this neuron is computed using a max-pooling operation, defined as:

$$o_j = \max\{z_j^1, z_j^2, ... z_j^{(T-t+1)}\}$$

The max-pooling operation provides location invariance to the neuron, i.e., the neuron is able to detect the features in the text regardless of where it appears.

The final output of the Convolutional Layer is the concatenation of the output from its m neurons, denoted by:

$$O = [o_1, o_2, ... o_m]$$

This output is then passed to a fully connected layer consisting of a weight matrix $W \in \mathbb{R}^{m \times n}$ and a bias $g \in \mathbb{R}^n$, which computes the final representation of the input text as:

$$x = \alpha(W \times O + g)$$

2.2 The DeepCoNN model

To compute the rating r_{AB} that $user_A$ would assign to $item_B$, the DeepCoNN model of [44] uses two *CNN Text Processors* side by side as shown in Figure 2. The first one processes the text labeled $text_A$, which consists of a concatenation of all the reviews that $user_A$ has written and produces a representation, x_A. Similarly, the second processes the text called $text_B$, which consists of a concatenation of all the reviews that have been written about $item_B$ and produces a representation, y_B. Both outputs are passed through a dropout layer [36]. Dropout is a function $\delta : \mathbb{R}^n \to \mathbb{R}^n$, that suppresses the

output of some of the neurons randomly and is a popular technique for regularizing a network. Let $\bar{x}_A = \delta(x_A)$ and $\bar{y}_B = \delta(y_B)$, denote the output of the dropout layer applied on x_A and y_B.

The model then concatenates the two representations as $z = [\bar{x}_A, \bar{y}_B]$ and passes it through a regression layer consisting of a *Factorization Machine* (FM) [32]. The FM computes the second order interactions between the elements of the input vector as:

$$\hat{r}_{AB} = w_0 + \sum_{i=1}^{|z|} w_i z_i + \sum_{i=1}^{|z|} \sum_{j=i+1}^{|z|} \langle \mathbf{v}_i, \mathbf{v}_j \rangle z_i z_j$$

where $w_0 \in \mathbb{R}$ is the global bias, $\mathbf{w} \in \mathbb{R}^{2n}$ weights each dimension of the input, and $\mathbf{V} \in \mathbb{R}^{2n \times k}$ assigns a k dimensional vector to each dimension of the input so that the pair-wise interaction between two dimensions i and j can be weighted using the inner product of the corresponding vectors $\mathbf{v}_i$ and $\mathbf{v}_j$. Note that the FM factorizes the pair-wise interaction, and therefore requires only $O(nk)$ parameters instead of $O(n^2)$ parameters which would have been required otherwise, where k is usually chosen such that $k \ll n$. This has been shown to give better parameter estimates under sparsity [32]. FMs have been used successfully in large scale recommendation services like online news[42].

FMs can be trained using different kinds of loss functions including least squared error (L_2), least absolute deviation (L_1), hinge loss and logit loss. In our experiments, L_1 loss gave a slightly better performance than L_2. DeepCoNN [44] also uses L_1 loss. Therefore, in this paper, all FMs are trained using L_1 loss, defined as:

$$loss = \sum_{(u_A, i_B, r_{AB}) \in D} |r_{AB} - \hat{r}_{AB}|$$

2.3 Limitations of DeepCoNN

DeepCoNN model has achieved impressive MSE values surpassing that of the previous state-of-the-art models that use review texts, like the Hidden Factors as Topics (HFT) model [24], Collaborative Topic Regression (CTR) [39] and Collaborative Deep Learning (CDL) [40], as well as Collaborative Filtering techniques that use only the rating information like Matrix Factorization (MF) [18] and Probabilistic Matrix Factorization (PMF) [33].

However, it was observed in our experiments that DeepCoNN achieves its best performance only when the text of the review written by the target user for the target item is available at test time. In real world recommendation settings, an item is always recommended to a user **before** they have experienced it. Therefore, it would be unreasonable to assume that the target review would be available at the time of testing.

Let rev_{AB} denote the review written by $user_A$ for an $item_B$. At training time, the text corresponding to $user_A$, denoted as $text_A$, consists of a concatenation of all reviews written by her in the training set. Similarly, the text for $item_B$, denoted by $text_B$, is a concatenation of all reviews written for that item in the training set. Both $text_A$ and $text_B$ includes rev_{AB} for all ($user_A, item_B$) pairs in the training set. At test time, there are two options for constructing the test inputs. For a test pair ($user_P, item_Q$), their pairwise review, rev_{PQ} in the test set, could be included in the texts corresponding to the user, $text_P$, and the item, $text_Q$, or could be omitted. In one of our datasets, the MSE obtained by DeepCoNN if rev_{PQ} is included in the test inputs is only 1.21. However, if rev_{PQ} is omitted, then the performance degrades severely to 1.89. This is lower than Matrix Factorization applied to the same dataset, which has an MSE of 1.86.
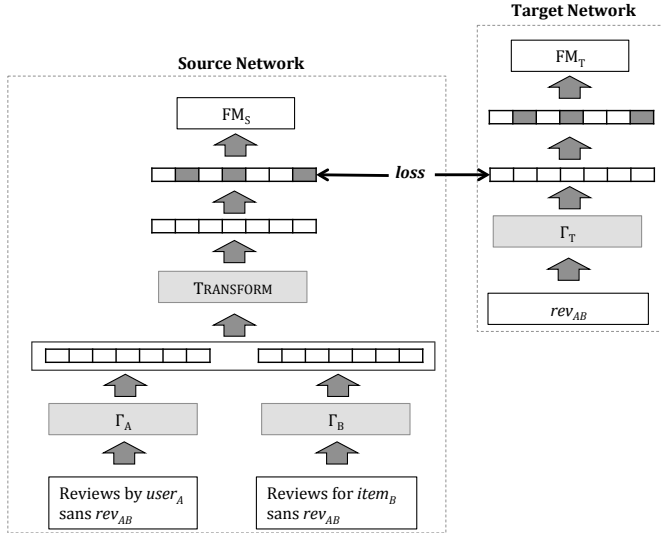
Figure 3: The TransNet architecture

If we train DeepCoNN in the setting that mimics the test setup, by omitting rev_{AB} in the texts of all $(user_A, item_B)$ pairs in the training set, the performance is better at 1.70, but still much higher than when rev_{AB} is available in both training and testing.

In the setting used in this paper, reviews in the validation and the test set are never accessed at any time, i.e., assumed to be unavailable — both during training and testing — simulating a real world situation.

2.4 TransNets

As we saw in the case of DeepCoNN, learning using the target review rev_{AB} at train time inadvertently makes the model dependent on the presence of such reviews at test time, which is unrealistic. However, as shown by the experiment above, rev_{AB} gives an insight into what $user_A$ thought about their experience with $item_B$, and can be an important predictor of the rating r_{AB}. Although unavailable at test time, rev_{AB} is available during training.

TransNet consists of two networks as shown in the architecture diagram of Figure 3, a *Target Network* that processes the target review rev_{AB} and a *Source Network* that processes the texts of the $(user_A, item_B)$ pair that does not include the joint review, rev_{AB}. Given a review text rev_{AB}, the Target Network uses a CNN Text Processor, Γ_T, and a Factorization Machine, FM_T, to predict the rating as:

$$x_T = \Gamma_T(rev_{AB})$$
$$\bar{x}_T = \delta(x_T)$$
$$\hat{r}_T = FM_T(\bar{x}_T)$$

Since the Target Network uses the actual review, its task is similar to sentiment analysis [19, 35].

The Source Network is like the DeepCoNN model with two CNN Text Processors, Γ_A for user text, *$text_A - rev_{AB}$*, and Γ_B for item text, *$text_B - rev_{AB}$*, and a Factorization Machine, FM_S, but with an additional TRANSFORM layer. The goal of the TRANSFORM layer is to **transform** the user and the item texts into an approximation of rev_{AB}, denoted by $r\hat{e}v_{AB}$, which is then used for predicting the rating. The Source Network predicts the rating as given below:

First, it converts the input texts into their latent form as:

$$x_A = \Gamma_A(text_A - rev_{AB})$$
$$x_B = \Gamma_B(text_B - rev_{AB})$$
$$z_0 = [x_A x_B]$$

The last step above is a concatenation of the two latent representations. This is then input to the TRANSFORM sub-network, which is a *L-layer deep non-linear transformational network*. Each layer l in TRANSFORM has a weight matrix $G_l \in \mathbb{R}^{n \times n}$ and bias $g_l \in \mathbb{R}^n$, and transforms its input z_{l-1} as:

$$z_l = \sigma(z_{l-1} G_l + g_l)$$

where σ is a non-linear activation function. Since the input to the first layer, z_0, is a concatenation of two vectors each of n dimensions, the first layer of TRANSFORM uses a weight matrix $G_1 \in \mathbb{R}^{2n \times n}$.

The output of the L^{th} layer of TRANSFORM, z_L is the approximation constructed by the TransNet for rev_{AB}, denoted by $r\hat{e}v_{AB}$. Note that we do not have to generate the surface form of rev_{AB}; It suffices to approximate $\Gamma_T(rev_{AB})$, the latent representation of the target review. The Source Network then uses this representation to predict the rating as:

$$\bar{z}_L = \delta(z_L)$$
$$\hat{r}_S = FM_S(\bar{z}_L)$$

During training, we will force the Source Network's representation z_L to be similar to the encoding of rev_{AB} produced by the Target Network, as we discuss below.

2.5 Training TransNets

TransNet is trained using 3 sub-steps as shown in Algorithm 1. In the first sub-step, for each training example (or a batch of such examples), the parameters of the Target Network, denoted by θ_T, which includes those of Γ_T and FM_T, are updated to minimize a L_1 loss computed between the actual rating r_{AB} and the rating $\hat{r}_T$ predicted from the actual review text rev_{AB}.

To teach the Source Network how to generate an approximation of the latent representation of the original review rev_{AB} generated by the Target Network, in the second sub-step, its parameters, denoted by θ_{trans}, are updated to minimize a L_2 loss computed between the transformed representation, $\bar{z}_L$, of the texts of the user and the item, and the representation x_T of the actual review. θ_{trans} includes the parameters of Γ_A and Γ_B, as well as the weights W_l and biases g_l in each of the transformation layers. θ_{trans} does not include the parameters of FM_S.

In the final sub-step, the remaining parameters of the Source Network, θ_S, which consists of the parameters of the FM_S are updated to minimize a L_1 loss computed between the actual rating r_{AB} and the rating $\hat{r}_S$ predicted from the transformed representation, $\bar{z}_L$. Note that each sub-step is repeated for each training example (or a batch of such examples), and not trained to convergence independently. The training method is detailed in Algorithm 1.

At test time, TransNet uses only the Source Network to make the prediction as shown in Algorithm 3.

2.6 Design Decisions and Other Architectural Choices

In this section, we describe some of the choices we have in designing the TransNet architecture and why they did not give good results in our preliminary experiments.

Algorithm 1 Training TransNet

1: **procedure** TRAIN(D_{train})
2:    **while** not converged **do**
3:       **for** $(text_A, text_B, rev_{AB}, r_{AB}) \in D_{train}$ **do**
4:          *#Step 1: Train Target Network on the actual review*
5:          $x_T = \Gamma_T(rev_{AB})$
6:          $\hat{r}_T = FM_T(\delta(x_T))$
7:          $loss_T = |r_{AB} - \hat{r}_T|$
8:          update θ_T to minimize $loss_T$
9:          *#Step 2: Learn to Transform*
10:         $x_A = \Gamma_A(text_A)$
11:         $x_B = \Gamma_B(text_B)$
12:         $z_0 = [x_A x_B]$
13:         $z_L = $ TRANSFORM(z_0)
14:         $\bar{z}_L = \delta(z_L)$
15:         $loss_{trans} = ||\bar{z}_L - x_T||_2$
16:         update θ_{trans} to minimize $loss_{trans}$
17:         *#Step 3: Train a predictor on the transformed input*
18:         $\hat{r}_S = FM_S(\bar{z}_L)$
19:         $loss_S = |r_{AB} - \hat{r}_S|$
20:         update θ_S to minimize $loss_S$
21:    **return** θ_{trans}, θ_S

Algorithm 2 Transform the input

1: **procedure** TRANSFORM(z_0)
2:    **for** layer $l \in L$ **do**
3:       $z_l = \sigma(z_{l-1} G_l + g_l)$
4:   **return** z_L

Algorithm 3 Testing using TransNet

1: **procedure** TEST(D_{test})
2:    **for** $(text_P, text_Q) \in D_{test}$ **do**
3:       *#Step 1: Transform the input*
4:       $x_P = \Gamma_A(text_P)$
5:       $x_Q = \Gamma_B(text_Q)$
6:       $z_0 = [x_P x_Q]$
7:       $z_L = $ TRANSFORM(z_0)
8:       $\bar{z}_L = \delta(z_L)$
9:       *#Step 2: Predict using the transformed input*
10:     $\hat{r}_{PQ} = FM_S(\bar{z}_L)$

2.6.1 ***Training with sub-steps vs. jointly.*** While training TransNets using Algorithm 1, in each iteration (or batch), we could choose to jointly minimize a total loss, $loss_{total} = loss_T + loss_{trans} + loss_S$. However, doing so will result in parameter updates to the target network, Γ_T, resulting from $loss_{trans}$, in addition to those from $loss_T$, i.e., the Target Network will get penalized for producing a representation that is different from that produced by the Source Network. This results in both networks learning to produce sub-optimal representations and converging to a lower performance in our experiments. Therefore, it is important to separate the Target Network's parameter updates so that it learns to produce the best representation which will enable it to make the most accurate rating predictions from the review text.
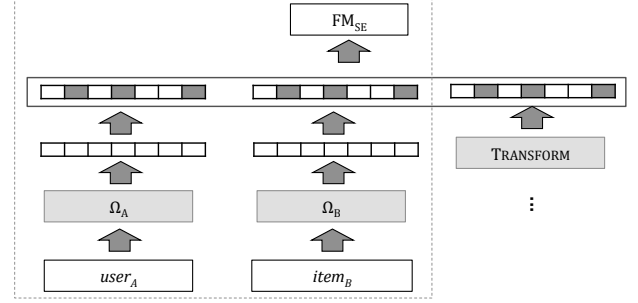
Figure 4: The Extended TransNet sub-architecture

2.6.2 ***Training Target Network to convergence independently.*** We could choose to first train the Target Network to convergence and then train the Source Network to emulate the trained Target Network. However, note that the Target Network's input is the actual review, which is unavailable for testing its performance, i.e., we do not know when the Target Network has converged with good generalization vs. when it is overfitting. The only way to measure the performance at test time is to check the output of the Source Network. Therefore, we let the Source and the Target Networks learn simultaneously and stop when the Source Network's test performance is good.

2.6.3 ***Using the same convolutional model to process text in both the Source and Target networks.*** We could choose to use the Γ_T that was trained in the Target Network to generate features from the user and the item text in the Source Network, instead of learning separate Γ_A and Γ_B. After all, we are learning to transform the latter's output into the former. However, in that case, TransNet would be constrained to generate generic features similar to topics. By providing it with separate feature generators, it can possibly learn to transform the occurrence of different features in the user text and the item text of the Source Network to another feature in the Target Network. For example, it could learn to transform the occurrence of features corresponding to say, 'love indian cuisine' & 'dislike long wait' in the user profile, and 'lousy service' & 'terrible chicken curry' in the item (restaurant) profile, to a feature corresponding to say, 'disappointed' in the target review, and subsequently predict a lower rating. Having separate feature generators in the Source Network gives TransNets more expressive power and gave a better performance compared to an architecture that reuses the Target Network's feature generator.

2.6.4 ***Training the TRANSFORM without the dropout.*** We could choose to match the output z_L of TRANSFORM with x_T instead of its dropped out version $\bar{z}_L$ in Step 2 of Algorithm 1. However, this makes the TRANSFORM layer unregularized, leading it to overfit thus giving poor performance.

2.7 Extended TransNets

TransNet uses only the text of the reviews and is user/item identity-agnostic, i.e., the user and the item are fully represented using the review texts, and their identities are not used in the model. However, in most real world settings, the identities of the users and items are known to the recommender system. In such a scenario, it is beneficial to learn a latent representation of the users and items, similar to Matrix Factorization methods. The *Extended TransNet* (TransNet-Ext) model achieves that by extending the architecture

Table 1: Dataset Statistics

Dataset	Category	#Users	#Items	#Ratings & Reviews
Yelp17		1,029,432	144,072	4,153,150
AZ-Elec	Electronics	4,200,520	475,910	7,820,765 (7,824,482)
AZ-CSJ	Clothing, Shoes and Jewelry	3,116,944	1,135,948	5,748,260 (5,748,920)
AZ-Mov	Movies and TV	2,088,428	200,915	4,606,671 (4,607,047)

of TransNet as shown in Figure 4.

The Source Network now has two embedding matrices Ω_A for users and Ω_B for items, which are functions of the form, $\Omega : id \rightarrow \mathbb{R}^n$. These map the string representing the identity of $user_A$ and $item_B$ into a n-dimensional representation. These latent representations are then passed through a dropout layer and concatenated with the output of the TRANSFORM layer before being passed to the FM regression layer. Therefore, given $user_A$ and $item_B$, TransNet-Ext computes the rating as:

$$\omega_A = \Omega(user_A)$$
$$\omega_B = \Omega(item_B)$$
$$\bar{z} = [\delta(\omega_A)\ \delta(\omega_B)\ \bar{z}_L]$$
$$\hat{r}_{SE} = FM_{SE}(\bar{z})$$

Computation of the loss in *Step 3* of Algorithm 1, $loss_{SE}$ is same as earlier: $loss_{SE} = |r_{AB} - \hat{r}_{SE}|$. But the parameter θ_S updated at the end now contains the embedding matrices Ω_A and Ω_B.

3 EXPERIMENTS AND RESULTS

3.1 Datasets

We evaluate the performance of the approach proposed in this paper on four large datasets. The first one, Yelp17, is from the latest Yelp dataset challenge[3], containing about 4M reviews and ratings of businesses by about 1M users. The rest are three of the larger datasets in the latest release of Amazon reviews[4] [25, 26] containing reviews and ratings given by users for products purchased on amazon.com, over the period of May 1996 - July 2014. We use the aggressively de-duplicated version of the dataset and also discard entries where the review text is empty. The statistics of the datasets are given in Table 1. The original size of the dataset before discarding empty reviews is given in brackets when applicable.

3.2 Evaluation Procedure and Settings

Each dataset is split randomly into train, validation and test sets in the ratio 80 : 10 : 10. After training on every 1000 batches of 500 training examples each, MSE is calculated on the validation and the test datasets. We report the MSE obtained on the test dataset when the MSE on the validation dataset was the lowest, similar to [24]. All algorithms, including the competitive baselines, were implemented in Python using TensorFlow[5], an open source software library for numerical computation, and were trained/tested on NVIDIA

GeForce GTX TITAN X GPUs. Training TransNet on Yelp17 takes approximately 40 minutes for 1 epoch (∼6600 batches) on 1 GPU, and gives the best performance in about 2–3 epochs.

Below are the details of the text processing and the parameter settings used in the experiments:

3.2.1 ***Text Pre-Processing and Embedding***. All reviews are first passed through a Stanford Core NLP Tokenizer [23] to obtain the tokens, which are then lowercased. Stopwords (the, and, is etc.) as well as punctuations are considered as separate tokens and are retained. A 64-dimensional $word2vec$[6] [27] embedding using the Skip-gram model is pre-trained on the 50,000 most frequent tokens in each of the training corpora.

3.2.2 ***CNN Text Processor***. We reuse most of the hyper-parameter settings reported by the authors of DeepCoNN [44] since varying them did not give any perceivable improvement. In all of the CNN Text Processors Γ_A, Γ_B and Γ_T, the number of neurons, m, in the convolutional layer is 100, the window size t is 3, and n, the dimension of the output of the CNN Text Processor, is 50. The maximum length of the input text, T, is set to 1000. If there are many reviews, they are randomly sorted and concatenated, and the first T tokens of the concatenated version are used. In our experiments, the word embedding dimension, d, is 64, and the vocabulary size, $|M|$ is 50,000. Also, the non-linearity, α, is $tanh$.

3.2.3 ***Dropout Layer and Factorization Machines***. All dropout layers have a keep probability of 0.5. In all of the factorization machines, FM_T, FM_S and FM_{SE}, the pair-wise interaction is factorized using a $k = 8$ dimensional matrix, V. Since FM_T processes a n-dimensional input, its parameters are $\mathbf{w}_T \in \mathbb{R}^n$ and $\mathbf{V}_T \in \mathbb{R}^{n \times k}$. Similarly, since FM_{SE} processes a $3n$-dimensional input, its parameters are $\mathbf{w}_{SE} \in \mathbb{R}^{3n}$ and $\mathbf{V}_{SE} \in \mathbb{R}^{3n \times k}$. All $\mathbf{w}$'s are initialized to 0.001, and all $\mathbf{V}$'s are initialized from a truncated normal distribution with 0.0 mean and 0.001 standard deviation. All FMs are trained to minimize an L_1 loss.

3.2.4 **TRANSFORM**. The default setting for the number of layers, L, is 2. We show the performance for different values of L in Section 3.5. All weight matrices G_l are initialized from a truncated normal distribution with 0.0 mean and 0.1 standard deviation, and all biases g_l are initialized to 0.1. The non-linearity, σ, is $tanh$.

3.2.5 ***TransNet-Ext***. The user (item) embedding matrices, Ω, are initialized from a random uniform distribution (-1.0, 1.0), and map users (items) that appear in the training set to a $n = 50$ dimensional space. New users (items) in the validation and test sets are mapped to a random vector.

3.2.6 ***Training***. All optimizations are learned using Adam [17], a stochastic gradient-based optimizer with adaptive estimates, at a learning rate set to 0.002. All gradients are computed by automatic differentiation in TensorFlow.

3.3 Competitive Baselines

We compare our method against the current state-of-the-art, DeepCoNN [44]. Since DeepCoNN was extensively evaluated against the previous state-of-the-art models like Hidden Factors as Topics (HFT) model [24], Collaborative Topic Regression (CTR) [39], Collaborative Deep Learning (CDL) [40] and Probabilistic Matrix Factorization (PMF) [33], and shown to surpass their performance by a wide margin, we refrain from repeating those comparisons in this paper. However, we do consider some variations of DeepCoNN.

Our competitive baselines are:

(1) **DeepCoNN**: The model proposed in [44]. During training, $text_A$ and $text_B$ corresponding to the $user_A$-$item_B$ pair contains their joint review rev_{AB}, along with reviews that $user_A$ wrote for other items and what other users wrote for $item_B$ in the training set. During testing, for a $user_P$-$item_Q$ pair, $text_P$ and $text_Q$ are constructed from only the training set and therefore, does not contain their joint review rev_{PQ}.

(2) **DeepCoNN-rev$_{AB}$**: The same DeepCoNN model (1) above, but trained in a setting that mimics the test setup, i.e., during training, $text_A$ and $text_B$ corresponding to the $user_A$-$item_B$ pair **does not contain** their joint review rev_{AB}, but only the reviews that $user_A$ wrote for other items and what other users wrote for $item_B$ in the training set. Testing procedure is the same as above: for a $user_P$-$item_Q$ pair, $text_P$ and $text_Q$ are constructed from only the training set and therefore, does not contain their joint review rev_{PQ} which is present in the test set.

(3) **MF**: A neural net implementation of Matrix Factorization with $n = 50$ latent dimensions.

We also provide the performance numbers of DeepCoNN in the setting where the test reviews are available at the time of testing. i.e. the same DeepCoNN model (1) above, but with the exception that at test time, for a $user_P$-$item_Q$ pair, $text_P$ and $text_Q$ are constructed from the training set as well as the test set, and therefore, contains their joint review rev_{PQ} from the test set. This is denoted as **DeepCoNN + Test Reviews**, and its performance is provided for the sole purpose of illustrating how much better the algorithm could perform, had it been given access to the test reviews.

3.4 Evaluation on Rating Prediction

Like prior work, we use the Mean Square Error (MSE) metric to evaluate the performance of the algorithms. Let N be the total number of datapoints being tested. Then MSE is defined as:

$$MSE = \frac{1}{N} \sum_{i=1}^{N} (r_i - \hat{r}_i)^2$$

where, r_i is the ground truth rating and $\hat{r}_i$ is the predicted rating for the i^{th} datapoint. Lower MSE indicates better performance.

The MSE values of the various competitive baselines are given in Table 2. For each dataset, the best score is highlighted in blue .

As can be seen from the Table, it is clear that TransNet and its variant TransNet-Ext perform better at rating prediction compared to the competitive baselines on all the datasets (p-value ≤ 0.05). It can also be seen that learning a user and item embedding using only the ratings in addition to the text helps TransNet-Ext improve the performance over the vanilla TransNet (p-value ≤ 0.1), except in the case of one dataset (`AZ-CSJ`).

It is also interesting to note that training DeepCoNN mimicking the test setup (DeepCoNN-rev$_{AB}$) gives a large improvement in the case of `Yelp`, but does not help in the case of the `AZ` datasets.

3.5 Picking the number of TRANSFORM layers

The TRANSFORM network uses L fully connected layers. In Figure 5, we plot the MSE of TransNet on the `Yelp17` dataset when varying L from 1 to 10. It can be seen from the figure that TransNets are quite robust to the choice of L, fluctuating only narrowly in its performance. Using only one layer gives the highest MSE, most


Figure 5: Variation in MSE with different Layers: TransNets on Yelp dataset

probably because it doesn't have enough parameters to learn how to transform the input. Using 10 layers also gives a high MSE, probably because it overfits or because it has too many parameters to learn. From the figure, using 2 or 5 layers gives the best MSE for this particular setting of TransNets. It is known that a 2 layer non-linear neural network is sufficient to represent all the logic gates including the XOR [11]. So, using 2 layers seems like a reasonable choice.

3.6 Finding the most helpful reviews

Our primary evaluation of TransNet is quantitative, using MSE of predicted ratings. We would also like to investigate whether the learned representation is qualitatively useful—i.e., does it capture interesting high-level properties of the user's review. One possible use of the learning representation would be to give the user information about her predicted reaction to the item that is more detailed than a rating. In this section, we show how TransNets could be used to find the most helpful reviews, personalized to each user. For example, the most helpful review for a user who is more concerned about the quality of service and wait times would be different from the most helpful review for another user who is sensitive to the price. For a test $user_P - item_Q$ pair, we run the Source Network with the text of their reviews from the training set to construct z_L, which is an approximation for the representation of their actual joint review. Candidate reviews are all the reviews rev_{CQ} in the training set written for $item_Q$ by other users. We pass each of them separately through the Target Network to obtain their latent representation $x_{CQ} = \Gamma_T(rev_{CQ})$. If rev_{CQ} had been most similar to what $user_P$ would write for $item_Q$, then x_{CQ} would be most similar to z_L. Therefore, to find the most helpful review, we simply choose the rev_{CQ} whose x_{CQ} is closest to z_L in Euclidean distance.

Some examples of such predicted most helpful reviews on the `Yelp17` dataset are listed in Table 3. Here, the column Original Review is the actual review that $user_P$ wrote for $item_Q$, and the column Predicted Review gives the most helpful of the candidate reviews predicted by TransNet. The examples show how the predicted reviews talk about particulars that the original reviews also highlight.

4 RELATED WORK

4.1 Recommendation Models

4.1.1 **Non-Neural Models**. The *Hidden Factors as Topics* (HFT) model [24] aims to find topics in the review text that are correlated

Table 2: Performance comparison using MSE metric

Dataset	DeepCoNN + Test Reviews	MF	DeepCoNN	DeepCoNN-rev$_{AB}$	TransNet	TransNet-Ext
Yelp17	1.2106	1.8661	1.8984	1.7045	1.6387	1.5913
AZ-Elec	0.9791	1.8898	1.9704	2.0774	1.8380	1.7781
AZ-CSJ	0.7747	1.5212	1.5487	1.7044	1.4487	1.4780
AZ-Mov	0.9392	1.4324	1.3611	1.5276	1.3599	1.2691

Table 3: Original review vs. Predicted most helpful review

Original Review	Predicted Review
my laptop flat lined and i did n't know why , just one day it did n't turn on . i cam here based on the yelp reviews and happy i did . although my laptop could n't be revived due to the fried motherboard , they did give me a full explanation about what they found and my best options . i was grateful they did n't charge me for looking into the problem , other places would have . i will definitely be coming back if needed . .	my hard drive crashed and i had to buy a new computer . the store where i bought my computer could n't get data off my old hard drive . neither could a tech friend of mine . works could ! they did n't charge for the diagnosis and only charged $ 100 for the transfer . very happy .
excellent quality korean restaurant . it 's a tiny place but never too busy , and quite possibly the best korean dumplings i 've had to date .	for those who live near by islington station you must visit this new korean restaurant that just opened up . the food too good to explain . i will just say i havent had a chance to take picture since the food was too grat .
this place is so cool . the outdoor area is n't as big as the fillmore location , but they make up for it with live music . i really like the atmosphere and the food is pretty spot on . the sweet potato fry dip is really something special . the vig was highly recommended to me , and i 'm passing that recommendation on to all who read this .	like going on monday 's . happy hour for drinks and apps then at 6pm their burger special . sundays are cool too , when they have live music on their patio .
i have attempted at coming here before but i have never been able to make it in because it 's always so packed with people wanting to eat . i finally came here at a good time around 6ish ... and not packed but by the time i left , it was packed ! the miso ramen was delicious . you can choose from add on 's on your soup but they charge you , i dont think they should , they should just treat them as condiments . at other ramen places that i have been too i get the egg , bamboo shoot , fire ball add on 's free . so i am not sure what their deal is .	hands down top three ramen spots on the west coast , right up there with , and the line can be just as long .
this place can be a zoo !! however , with the produce they have , at the prices they sell it at , it is worth the hassle . be prepared to be pushed and shoved . this is much the same as in asia . my wife (from vietnam) says that the markets in asia are even more crowded . i agree as i have seen vietnam with my own eyes .	i enjoy going to this market on main street when i am ready to can ... the prices are great esp for onions . . broccoli and bell peppers ... a few times they have had bananas for $ 3.00 for a huge box like 30 lbs ... you can freeze them or cover in ... or make banana bread if they begin to go dark ... and ripe . the employees will talk if you say hello first ...
great spot for outdoor seating in the summer since it 's sheltered early from the sun . good service but americanos sometimes are not made right	this is my " go to " starbucks due to the location being close to where i live . i normally go through the drive-thru , which more likely than not , has a long line . . but does n't every starbucks ? i have always received great customer service at this location ! there has been a couple times that my order was n't correct - which is frustrating when you are short on time & depend on your morning coffee ! but overall you should have a good experience whether you drive-thru or dine in !
this is my favorite place to eat in south charlotte . great cajun food . my favorite is the fried oysters with cuke salad and their awesome mac 'n' cheese (their mac 'n' cheese is not out of a box) . their sweet tea would make my southern grandma proud . to avoid crowds , for lunch i recommend arriving before 11:30 a.m. or after 1 p.m. and for dinner try to get there before 6 p.m. is not open on sundays .	always !! too small location so wait line can be long . been going to for 13 years .
very overpriced food , very substandard food . wait staff is a joke . if you like being ignored in the front room of a one story house and being charged for it , by all means . otherwise , go to freaking mcdonald 's .	i want this place to be good but it fails me every time . i brought my parents here for dinner and was totally embarrassed with my dining choice . i tried it two more times after that and continue to be disappointed . their menu looks great but what is delivered is a total let down . to top it off , the service is painfully slow . the only thing this place has going for it is the dog friendly patio and craft beers . i hope someone reads these reviews as the poor service piece continues to be brought up as an issue .

with the latent parameters of users. They propose a transformation function which converts user's latent factors to the topic distribution of the review, and since the former exactly defines the latter, only one of them is learned. A modified version of HFT is the *TopicMF* model [4], where the goal is to match the latent factors learned for the users and items using MF with the topics learned on their joint reviews using a Non-Negative Matrix Factorization, which is then jointly optimized with the rating prediction. In their transformation function, the proportion of a particular topic in the review is a linear combination of its proportion in the latent factors of the user and the item, which is then converted into a probability distribution over all topics in that review. Unlike these two models, TransNet computes each factor in the transformed review from a non-linear combination of any number of factors from the latent representations of either the user or the item or both. Another extension to HFT is the *Rating Meets Reviews* (RMR)

model [22] where the rating is sampled from a Gaussian mixture.

The *Collaborative Topic Regression* (CTR) model proposed in [39] is a content based approach, as opposed to a context / review based approach. It uses LDA [5] to model the text of documents (scientific articles), and a combination of MF and content based model for recommendation. The *Rating-boosted Latent Topics* (RBLT) model of [37] uses a simple technique of repeating a review r times in the corpus if it was rated r, so that features in higher rated reviews will dominate the topics. *Explicit Factor Models* (EFM) proposed in [43] aims to generate explainable recommendations by extracting explicit product features (aspects) and users' sentiments towards these aspects using phrase-level sentiment analysis.

4.1.2 Neural Net Models. The most recent model to successfully employ neural networks at scale for rating prediction is the *Deep Cooperative Neural Networks* (DeepCoNN) [44], which was discussed in detail in Section 2. Prior to that work, [2] proposed two models: *Bag-of-Words regularized Latent Factor* model (BoWLF) and *Language Model regularized Latent Factor* model (LMLF), where MF was used to learn the latent factors of users and items, and likelihood of the review text, represented either as a bag-of-words or an LSTM embedding [14], was computed using the item factors. [34] proposed a CNN based model identical to DeepCoNN, but with attention mechanism to construct the latent representations, the inner product of which gave the predicted ratings.

Some of the other past research uses neural networks in a CF setting with content, but not reviews. The *Collaborative Deep Learning* (CDL) model [40] uses a Stacked De-noising Auto Encoder (SDAE) [38] to learn robust latent representations of items from their content, which is then fed into a CTR model [39] for predicting the ratings. A very similar approach to CDL is the *Deep Collaborative Filtering* (DCF) method [21] which uses Marginalized De-noising Auto-Encoder (mDA) [7] instead. The *Convolutional Matrix Factorization* (ConvMF) model [16] uses a CNN to process the description associated with the item and feed the resulting latent vectors into a PMF model for rating prediction. The *Multi-View Deep Neural Net* (MV-DNN) model [10] uses a deep neural net to map user's and item's content into a shared latent space such that their similarity in that space is maximized. [29] proposed to generate the latent factors of items – music in this case— from the content, audio signals. The predicted latent factors of the item were then used in a CF style with the latent factors of the user. [3] also proposed a similar technique but adapted to recommending scientific-articles. [9] used a deep neural net to learn a latent representation from video content which is then fed into a deep ranking network.

Prior research has also used deep neural nets for learning latent factors from ratings alone, i.e., without using any content or review. *Collaborative De-noising Auto-Encoder* model (CDAE) [41] learns to reconstruct user's feedback from a corrupted version of the same.

4.2 Comparison to Related Architectures

4.2.1 Student-Teacher Models. Student-Teacher models [6, 13] also have two networks: a Teacher Network, which is large and complex, and typically an ensemble of different models, is first trained to make predictions, and a much simpler Student Network, which learns to emulate the output of the Teacher Network, is trained later. There are substantial differences between Student-Teacher models and TransNets in how they are structured. Firstly, in Student-Teacher models, the input to both the student and the

teacher models are the same. For example, in the case of digit recognition, both networks input the same image of the digit. However, in TransNets, the inputs to the two networks are different. In the Target, there is only one input – the review by $user_A$ for an $item_B$ designated as rev_{AB}. But, in the Source, there are two inputs: all the reviews written by $user_A$ sans rev_{AB} and all the reviews written for $item_B$ sans rev_{AB}. Secondly, in Student-Teacher models, the Teacher is considerably complex in terms of width and depth, and the Student is more light-weight, trying to mimic the Teacher. In TransNets, the complexities are reversed. The Target is lean while the Source is heavy-weight, often processing large pieces of text using twice the number of parameters as the Target. Thirdly, in Student-Teacher models, the Teacher is pre-trained whereas in TransNets the Target is trained simultaneously with the Source. A recently proposed Student-Teacher model in [15] does train both the Student and the Teacher simultaneously. Also, in Student-Teacher models, the emphasis is on learning a simpler and easier model that can achieve similar results as a very complex model. But in TransNets, the objective is to learn how to transform a source representation to a target representation.

4.2.2 Generative Adversarial Networks. TransNets also bear semblance to GANs [12, 31] since both are attempting to generate an output which is similar to realistic data. But the models are fundamentally different. Firstly, unlike GAN where the Generative network generates an output from a random starting point, TransNets have a starting point for each example – the reviews written by the user and those written for the item. Secondly, the Adversarial network in GAN tries to classify if the output is real or synthetic. In TransNets, although the objective is to minimize the dissimilarity between the generated representation and that of the real input, there is no adversarial classifier that attempts to separate each out. Thirdly, in GANs, the adversarial network needs to learn a notion of 'real' outputs, which is quite generic. In TransNets, there is always a specific real output to compare to and does not need to learn what a generic real output will look like.

5 CONCLUSIONS

Using reviews for improving recommender systems is an important task and is gaining a lot of attention in the recent years. A recent neural net model, *DeepCoNN*, uses the text of the reviews written by the user and for the item to learn their latent representations, which are then fed into a regression layer for rating prediction. However, its performance is dependent on having access to the user-item pairwise review, which is unavailable in real-world settings.

In this paper, we propose a new model called *TransNets* which extends *DeepCoNN* with an additional TRANSFORM layer. This additional layer learns to transform the latent representations of user and item into that of their pair-wise review so that at test time, an approximate representation of the target review can be generated and used for making the predictions. We also showed how *TransNets* can be extended to learn user and item representations from ratings only which can be used in addition to the generated review representation. Our experiments showed that *TransNets* and its extended version can improve the state-of-the-art substantially.

REFERENCES

[1] Gediminas Adomavicius and Alexander Tuzhilin. 2008. Context-aware Recommender Systems. In *Proceedings of the 2008 ACM Conference on Recommender Systems (RecSys '08)*. 335–336.

[2] Amjad Almahairi, Kyle Kastner, Kyunghyun Cho, and Aaron Courville. 2015. Learning Distributed Representations from Reviews for Collaborative Filtering. In *Proceedings of the 9th ACM Conference on Recommender Systems (RecSys '15)*. 147–154.

[3] Trapit Bansal, David Belanger, and Andrew McCallum. 2016. Ask the GRU: Multi-task Learning for Deep Text Recommendations. In *Proceedings of the 10th ACM Conference on Recommender Systems (RecSys '16)*. 107–114.

[4] Yang Bao, Hui Fang, and Jie Zhang. 2014. TopicMF: Simultaneously Exploiting Ratings and Reviews for Recommendation. In *Proceedings of the Twenty-Eighth AAAI Conference on Artificial Intelligence (AAAI'14)*. 2–8.

[5] David M. Blei, Andrew Y. Ng, and Michael I. Jordan. 2003. Latent Dirichlet Allocation. *J. Mach. Learn. Res.* 3 (March 2003), 993–1022.

[6] Cristian Buciluǎ, Rich Caruana, and Alexandru Niculescu-Mizil. 2006. Model Compression. In *Proceedings of the 12th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining (KDD '06)*. 535–541.

[7] Minmin Chen, Zhixiang Xu, Kilian Q. Weinberger, and Fei Sha. 2012. Marginalized Denoising Autoencoders for Domain Adaptation. In *Proceedings of the 29th International Coference on International Conference on Machine Learning (ICML'12)*. 1627–1634.

[8] Ronan Collobert, Jason Weston, Léon Bottou, Michael Karlen, Koray Kavukcuoglu, and Pavel Kuksa. 2011. Natural Language Processing (Almost) from Scratch. *J. Mach. Learn. Res.* 12 (Nov. 2011), 2493–2537.

[9] Paul Covington, Jay Adams, and Emre Sargin. 2016. Deep Neural Networks for YouTube Recommendations. In *Proceedings of the 10th ACM Conference on Recommender Systems (RecSys '16)*. 191–198.

[10] Ali Mamdouh Elkahky, Yang Song, and Xiaodong He. 2015. A Multi-View Deep Learning Approach for Cross Domain User Modeling in Recommendation Systems. In *Proceedings of the 24th International Conference on World Wide Web (WWW '15)*. 278–288.

[11] Ian Goodfellow, Yoshua Bengio, and Aaron Courville. 2016. *Deep Learning*. MIT Press. http://www.deeplearningbook.org.

[12] Ian J. Goodfellow, Jean Pouget-Abadie, Mehdi Mirza, Bing Xu, David Warde-Farley, Sherjil Ozair, Aaron Courville, and Yoshua Bengio. 2014. Generative Adversarial Nets. In *Proceedings of the 27th International Conference on Neural Information Processing Systems (NIPS'14)*. 2672–2680.

[13] Geoffrey E. Hinton, Oriol Vinyals, and Jeffrey Dean. 2014. Distilling the Knowledge in a Neural Network. In *Deep Learning and Representation Learning Workshop (NIPS '14)*.

[14] Sepp Hochreiter and Jürgen Schmidhuber. 1997. Long Short-Term Memory. *Neural Computation* 9, 8 (Nov. 1997), 1735–1780.

[15] Zhiting Hu, Xuezhe Ma, Zhengzhong Liu, Eduard Hovy, and Eric Xing. 2016. Harnessing Deep Neural Networks with Logic Rules. In *Proceedings of the 54th Annual Meeting of the Association for Computational Linguistics (ACL)*. 2410–2420.

[16] Donghyun Kim, Chanyoung Park, Jinoh Oh, Sungyoung Lee, and Hwanjo Yu. 2016. Convolutional Matrix Factorization for Document Context-Aware Recommendation. In *Proceedings of the 10th ACM Conference on Recommender Systems (RecSys '16)*. 233–240.

[17] Diederik P. Kingma and Jimmy Ba. 2014. Adam: A Method for Stochastic Optimization. In *Proceedings of the 3rd International Conference on Learning Representations (ICLR)*.

[18] Yehuda Koren. 2008. Factorization Meets the Neighborhood: A Multifaceted Collaborative Filtering Model. In *Proc. KDD '08*. 426–434.

[19] Quoc V Le and Tomas Mikolov. 2014. Distributed Representations of Sentences and Documents.. In *ICML*, Vol. 14. 1188–1196.

[20] Y. Lecun, L. Bottou, Y. Bengio, and P. Haffner. 1998. Gradient-based learning applied to document recognition. *Proc. IEEE* 86, 11 (1998), 2278–2324.

[21] Sheng Li, Jaya Kawale, and Yun Fu. 2015. Deep Collaborative Filtering via Marginalized Denoising Auto-encoder. In *Proceedings of the 24th ACM International on Conference on Information and Knowledge Management (CIKM '15)*. 811–820.

[22] Guang Ling, Michael R. Lyu, and Irwin King. 2014. Ratings Meet Reviews, a Combined Approach to Recommend. In *Proceedings of the 8th ACM Conference on Recommender Systems (RecSys '14)*. 105–112.

[23] Christopher D. Manning, Mihai Surdeanu, John Bauer, Jenny Finkel, Steven J. Bethard, and David McClosky. 2014. The Stanford CoreNLP Natural Language Processing Toolkit. In *Association for Computational Linguistics (ACL) System Demonstrations*. 55–60.

[24] Julian McAuley and Jure Leskovec. 2013. Hidden Factors and Hidden Topics: Understanding Rating Dimensions with Review Text. In *Proceedings of the 7th ACM Conference on Recommender Systems (RecSys '13)*. 165–172.

[25] Julian McAuley, Rahul Pandey, and Jure Leskovec. 2015. Inferring Networks of Substitutable and Complementary Products. In *Proceedings of the 21th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining (KDD '15)*. 785–794.

[26] Julian McAuley, Christopher Targett, Qinfeng Shi, and Anton van den Hengel. 2015. Image-Based Recommendations on Styles and Substitutes. In *Proceedings*

of the 38th International ACM SIGIR Conference on Research and Development in Information Retrieval (SIGIR '15). 43–52.

[27] Tomas Mikolov, Ilya Sutskever, Kai Chen, Greg Corrado, and Jeffrey Dean. 2013. Distributed Representations of Words and Phrases and Their Compositionality. In *Proceedings of the 26th International Conference on Neural Information Processing Systems (NIPS'13)*. 3111–3119.

[28] Vinod Nair and Geoffrey E. Hinton. 2010. Rectified Linear Units Improve Restricted Boltzmann Machines. In *Proceedings of the 27th International Conference on Machine Learning (ICML-10), June 21-24, 2010, Haifa, Israel*. 807–814.

[29] Aäron van den Oord, Sander Dieleman, and Benjamin Schrauwen. 2013. Deep Content-based Music Recommendation. In *Proceedings of the 26th International Conference on Neural Information Processing Systems (NIPS'13)*. 2643–2651.

[30] Jeffrey Pennington, Richard Socher, and Christopher D. Manning. 2014. GloVe: Global Vectors for Word Representation. In *Empirical Methods in Natural Language Processing (EMNLP)*. 1532–1543.

[31] Scott Reed, Zeynep Akata, Xinchen Yan, Lajanugen Logeswaran, Bernt Schiele, and Honglak Lee. 2016. Generative Adversarial Text to Image Synthesis. In *Proceedings of the 33rd International Conference on International Conference on Machine Learning - Volume 48 (ICML'16)*. 1060–1069.

[32] Steffen Rendle. 2010. Factorization Machines. In *Proceedings of the 2010 IEEE International Conference on Data Mining (ICDM '10)*. 995–1000.

[33] Ruslan Salakhutdinov and Andriy Mnih. 2007. Probabilistic Matrix Factorization. In *Proceedings of the 20th International Conference on Neural Information Processing Systems (NIPS'07)*. Curran Associates Inc., USA, 1257–1264. http://dl.acm.org/citation.cfm?id=2981562.2981720

[34] Sungyong Seo, Jing Huang, Hao Yang, and Yan Liu. 2017. Representation Learning of Users and Items for Review Rating Prediction Using Attention-based Convolutional Neural Network. In *3rd International Workshop on Machine Learning Methods for Recommender Systems (MLRec) (SDM '17)*.

[35] Richard Socher, Alex Perelygin, Jean Wu, Jason Chuang, Christopher D. Manning, Andrew Y. Ng, and Christopher Potts. 2013. Recursive Deep Models for Semantic Compositionality Over a Sentiment Treebank. In *Proceedings of the 2013 Conference on Empirical Methods in Natural Language Processing*. Association for Computational Linguistics, 1631–1642.

[36] Nitish Srivastava, Geoffrey Hinton, Alex Krizhevsky, Ilya Sutskever, and Ruslan Salakhutdinov. 2014. Dropout: A Simple Way to Prevent Neural Networks from Overfitting. *J. Mach. Learn. Res.* 15, 1 (2014), 1929–1958.

[37] Yunzhi Tan, Min Zhang, Yiqun Liu, and Shaoping Ma. 2016. Rating-boosted Latent Topics: Understanding Users and Items with Ratings and Reviews. In *Proceedings of the Twenty-Fifth International Joint Conference on Artificial Intelligence (IJCAI'16)*. 2640–2646.

[38] Pascal Vincent, Hugo Larochelle, Isabelle Lajoie, Yoshua Bengio, and Pierre-Antoine Manzagol. 2010. Stacked Denoising Autoencoders: Learning Useful Representations in a Deep Network with a Local Denoising Criterion. *J. Mach. Learn. Res.* 11 (Dec. 2010), 3371–3408.

[39] Chong Wang and David M. Blei. 2011. Collaborative Topic Modeling for Recommending Scientific Articles. In *Proceedings of the 17th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining (KDD '11)*. 448–456.

[40] Hao Wang, Naiyan Wang, and Dit-Yan Yeung. 2015. Collaborative Deep Learning for Recommender Systems. In *Proceedings of the 21th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining (KDD '15)*. 1235–1244.

[41] Yao Wu, Christopher DuBois, Alice X. Zheng, and Martin Ester. 2016. Collaborative Denoising Auto-Encoders for Top-N Recommender Systems. In *Proceedings of the Ninth ACM International Conference on Web Search and Data Mining (WSDM '16)*. 153–162.

[42] Xing Yi, Liangjie Hong, Erheng Zhong, Nanthan Nan Liu, and Suju Rajan. 2014. Beyond Clicks: Dwell Time for Personalization. In *Proceedings of the 8th ACM Conference on Recommender Systems (RecSys '14)*. 113–120.

[43] Yongfeng Zhang, Guokun Lai, Min Zhang, Yi Zhang, Yiqun Liu, and Shaoping Ma. 2014. Explicit Factor Models for Explainable Recommendation Based on Phrase-level Sentiment Analysis. In *Proceedings of the 37th International ACM SIGIR Conference on Research & Development in Information Retrieval (SIGIR '14)*. 83–92.

[44] Lei Zheng, Vahid Noroozi, and Philip S. Yu. 2017. Joint Deep Modeling of Users and Items Using Reviews for Recommendation. In *Proceedings of the Tenth ACM International Conference on Web Search and Data Mining (WSDM '17)*. 425–434.